Project 5: Fault-tolerant SurfStore

FAQ/Updates

Overview

In Project 4, you built a DropBox clone called "SurfStore". Because data blocks are immutable and cannot be updated (since doing so would change their hash values, and thus they'd become entirely new blocks), replicating blocks is quite easy. On the other hand, replicating the MetaStore service is quite challenging, because multiple clients can update the Metadata of a file in a concurrent manner. To ensure that the Metadata store is fault tolerant and stays consistent regardless of failures, we can implement it as a replicated state machine design, which is the purpose of this project.

In this project, you are going to modify your metadata server to make it fault tolerant based on the RAFT protocol. To make the project reasonable to complete in 2 weeks you will only implement the log replication part of the protocol. The exact differences are explained in this document.

You will re-use your P4 solution for the "starter code" for P5. Instructions on how to copy the files over from the P5 starter code are in the README. You **must** implement the interfaces in RaftInterfaces.go and your server **must** be defined as RaftSurfstoreServer. In P5, we're only testing the metadata functions (updatefile and getfileinfo), we will not test the putblock, getblock, etc. functions. These should still be functional however so that your client can successfully upload and download files. To help with testing we have released a few test cases, and you should implement your own as well.

Review

- [The RAFT paper](#)
  - Section 1, 2, 4, 5, 8, and 11 are required reading
  - Sections 3, 6, 7, 9, 10, and 12 are optional and not necessary for your project
  - You will **not** be implementing log compaction or membership changes in this project!
- [The RAFT website](#)
- [Very helpful visualization of the protocol](#)
- [The RAFT simulator](#)
  - If you have questions about what "should" happen during certain circumstances, this simulator is your best resource for investigating that situation.

Design

You will implement a RaftSurfstoreServer which functions as a fault tolerant MetaStore from project 4. Each RaftSurfstoreServer will communicate with other RaftSurfstoreServers via GRPC. Each server is aware of all other possible servers (from the configuration file), and new servers do not dynamically join the cluster (although existing servers can "crash" via the Crash api). Leaders will be set through the SetLeader API call, so there are no elections.

Using the protocol, if the leader can query a majority quorum of the nodes, it will reply back to the client with the correct answer. As long as a majority of the nodes are up and not in a crashed state, the clients should be able to interact with the system successfully. When a majority of nodes are in a crashed state, clients should block and not receive a response until a majority are restored. Any clients that interact with a non-leader should get an error message and retry to find the leader.

ChaosMonkey

To test your implementation you will need to use the RaftTestingInterface which defines 3 functions for 'chaos' testing and one function to access the internal state. This simulates the server crashing and failing. Note that we won't really crash your program (e.g. by typing "Control-C" or sending it the kill command). When the Crash() call is given to a server, the server should enter a crashed state, and if it gets any AppendEntries or other calls, it should reply back with an error and not update its internal state. If a client tries to contact a crashed node, it should just return an error indicating it is crashed (letting the client search for the actual leader).

Our autograding code will call the Crash/Restore/IsCrashed/GetInternalState methods as part of testing your codebase. To ensure your code works correctly, you should implement your own tests that invoke these methods to ensure that the properties of RAFT hold up in your solution.

API summary

As a quick summary of the calls you need in P5:

RPC call	Description	Who calls this?	Response during "crashed" state	GRPC Output
AppendEntries()	Replicates log entries; serves as a heartbeat mechanism	Your server code only	Should return an "isCrashed" error; procedure has no effect if server is crashed	AppendEntryOutput: should have the appropriate fields as described in the Raft paper

SetLeader()	Emulates elections, sets the node to be the leader	The autograder, your testing code	Should return an "isCrashed" error; procedure has no effect if server is crashed	Success: True if the server was successfully made the leader
SendHeartbeat()	Sends a round of AppendEntries to all other nodes. The leader will attempt to replicate logs to all other nodes when this is called. It can be called even when there are no entries to replicate. If a node is not in the leader state it should do nothing.	The autograder, your testing code	Should return an "isCrashed" error; procedure has no effect if server is crashed	We will not check the output of SendHeartbeat. You can return True always or have some custom logic based on your implementation.
getblock(), putblock(), hasblocks()	Procedures related to the contents of files	Your client	N/A	Same as P4
GetBlockStoreAddr()	Returns the block store address	Your client	If the node is the leader, and if a majority of the nodes are working, should return the correct answer; if a majority of the nodes are crashed, should block until a majority recover. If not the leader, should indicate an error back to the client	Same as P4
GetFileInfoMap()	Returns metadata from	Your client	If the node is the leader, and if a	Same as P4

	the filesystem		majority of the nodes are working, should return the correct answer; if a majority of the nodes are crashed, should block until a majority recover. If not the leader, should indicate an error back to the client	
UpdateFile()	Updates a file's metadata	Your client	If the node is the leader, and if a majority of the nodes are working, should return the correct answer; if a majority of the nodes are crashed, should block until a majority recover. If not the leader, should indicate an error back to the client	Same as P4
GetInternalState()	Returns the internal state of a Raft server	The autograder	Always return the state	RaftInternalState The correct output is implemented in the starter code, do not change
Crash()	Cause the server to enter a "crashed" state	The autograder; Your testing code	Crashing a server that is already crashed has no effect	Success: Value does not matter, if you want to change for your testing you can
Restore()	Causes the server to no longer be crashed	The autograder; Your testing code	Causes the server to recover and no longer be crashed	Success: Value does not matter, if you want to change

				for your testing you can
IsCrashed()	Returns whether the server is in a crashed state	The autograder, your testing code	Should return the correct state	CrashedState: True if the server is crashed, else false

Please consult Figure 2 in the Raft paper, and the SurfStore.proto and RaftInterfaces.go files for more information.

Changes to the command-line arguments of your server

Your server code needs to handle different command line arguments than in project 4. For project 5, your server should take in a path to a configuration file, the ID number of that server, and an address to a BlockStore server. For example:

```
$ go run cmd/SurfstoreRaftServerExec/main.go -f configfile.txt -i 0 -b
localhost:8080
```

Would start the server with a configuration file of configfile.txt. It would tell the server that it is server 0 in the configured servers. And the server that starts will assume that there is a BlockStore running on localhost:8080.

The client now also takes the same configuration file so that it knows where the servers are running. You can run the client with:
```
$ go run cmd/SurfstoreClientExec/main.go -f configfile.txt baseDir blockSize
```

A Makefile is provided to things easier, for example to run a BlockStore you should type:
```
$ make run-blockstore
```

Then in a separate terminal you can run a raft server with:
```
$ make IDX=0 run-raft
```

Configuration file

Your server will receive a configuration file as part of its initialization. The format is as follows:

```
M: 3
metadata0: <host>:<port>
metadata1: <host>:<port>
metadata2: <host>:<port>
```

As an example:

```
M: 5
metadata0: localhost:9001
metadata1: localhost:9002
metadata2: localhost:9003
metadata3: localhost:9004
metadata4: localhost:9005
```

The metadata numbers start at zero. Note that all of the configuration files and command line arguments we provide to your system will be legal and we won't introduce syntax errors or other errors in your configuration files. You should be able to handle a variable number of servers, though you don't need to handle more than 10. Make sure to test with an even number of servers.

Leader election

Elections do not take place; the leader is set deterministically through a SetLeader API. The SetLeader function should emulate an election, so after calling it on a node it should set all the state as if that node had just won an election. Because elections do not take place, there are only two node states, Leader and Follower. The candidate state does not exist. You can guarantee that a SendHeartbeat call will be issued after every call to SetLeader. There will not be any tests where these functions are called in such a way that the nodes enter a state that is impossible in the Raft protocol. Leader conflicts are handled on Heartbeat or AppendEntries call. For example:
(Node A).SetLeader()
(Node A).SendHeartbeat()
…. Here node A is the leader in term 1
(Node B).SetLeader()
… here node A and B both think they are the leader, however node B is in term 2 and node A is term 1
(Node B).SendHeartbeat()
… Node A gets a heartbeat message from node B which has a higher term, node A steps down

Heartbeat

There is no heartbeat timer. Heartbeats are triggered through a SendHeartbeat API. There is no heartbeat countdown, so once the node is in the leader state, it stays in that state until there is another leader. A node could find out about another leader either through receiving a heartbeat from another leader with a higher term, or receiving a RPC response that has a higher term.

Clients

The client will keep a list of Raft servers, when making a request if the client gets an ERR_NOT_LEADER error, it can simply try the next server in the list. The server does not need

to include the last known leader in its response. If the client is not able to successfully complete its GRPC request from any servers in the list it can return an error.

RPC limits

A single node should be able to handle up to O(10) RPC calls per second. Do not create an implementation with an excessive number of RPC calls (well above this limit), as that might cause our testing framework to malfunction.

Persistent storage

Because we're only simulating crashes, you do not need to persist your replicated log on the filesystem.

Testing

Several example tests are provided in test/raft_test.go and test/basic_test.go. Some example config files are in test/config_files and example files are in test/test_files. You should also create your own files and configs to test your code. Any test code you write will be overwritten before running the autograder.

Debug Logging

Before submitting your code, try to minimize excessive logging. If there are a ton of log messages, it can interfere with the grader. Please make sure that all log statements end with a newline. Our autograder script separates results by newlines and having a logging statement can interfere with the output.

Design notes

Committing entries

When a client sends a command to the leader, the leader is going to log that command in its local log, then issue a two-phase commit operation to its followers. When a majority of those followers approve of the update, the leader can commit the transaction locally, apply the log to its state machine, and send the success response to the client. If the leader failed to get approval from all of the followers (maybe because one or more are crashed), it should keep trying indefinitely. The commit behavior is described in the Raft paper, section 5.3

Versions and leaders

updatefile() should only be applied when the given version number is exactly one higher than the version stored in the leader. Every operation, including updatefile() and getfileinfomap() needs to involve talking to a majority of the nodes.

Note that the followers (and leaders) need to always implement GetInternalState, even when they are crashed.

SendHeartbeat

You are guaranteed to have SendHeartbeat called:
1. After every call to SetLeader the node that had SetLeader called will have SendHeartbeat called.
2. After every UpdateFile call the node that had UpdateFile called will have SendHeartbeat called.
3. After the test, the leader will have SendHeartbeat called one final time. Then all of the nodes should be ready for the internal state to be collected through GetInternalState.
4. After every SyncClient operation

Starter code

Follow the instructions in the starter code to extend your project 4 code. From our project 4 solution, we changed RaftSurfstoreServer.go, RaftUtils.go, and SurfstoreRPCClient.go. Our project 5 solution made ~600 insertions and ~100 deletions (from our project 4 solution).